

# Data cleaning

LPO 9951 | Fall 2020

## PURPOSE

Data cleaning is a term which refers to a variety of highly applied activities that are intended to ensure that the data are not, in some way, incorrect. Data cleaning is the least glamorous part of data analysis (if any part of data analysis can be called “glamorous”) and by far the most important. Time spent on this activity is never wasted.

A “clean” dataset is one in which every variable is labeled appropriately, and organized in a way that will make sense to a user. If missing data exists, there is a unified approach to dealing with it (i.e., every missing data point is coded in the same way). The data have been checked for consistency using a variety of more or less common sense approaches. It’s important to note from the outset that data-cleaning is not a specific data analytic approach so much as an attitude about the data, combined with some specific techniques. Here are some, but only some, of those techniques.

We’ll be using a deliberately “dirtied up” dataset. This is based on the `caschools` data, which contains data on all of the school districts in the state of California.

```
. capture log close                                // closes any logs, should they be open

. set linesize 90

. log using "cleaning.log", replace                // open new log
-----
      name: <unnamed>
      log:  /Users/doylewr/lpo_prac/lessons/s1-07-cleaning/cleaning.log
  log type: text
opened on:  14 Oct 2020, 11:14:59

. clear all                                         // clear memory

. set more off                                     // turn off annoying "__more__" feature
```

## Variable labels

Not everyone you work with will be as well-trained as you. Large datasets that are appended over time (think institutional data) also have a tendency to suffer from drift or entropy. Many times you will receive a dataset with incomplete or

incomprehensible variable labels. You will need to figure out what to do with these. The labels in this example dataset are obviously a mess.

```
. use caschool_problem, replace
```

## Techniques for working with single variables

Each of the following describes techniques that can help you to find data points that may be unreliable or wrong. When you find such a data point, you need to rely on your judgment and the context of the problem to decide what to do next.

### Outliers

Look for outliers: values that are extraordinarily far from the mean or median (e.g., more than 3 s.d. away for approximately normal data).

Let's take a look at student teacher ratio using the `boxplot` and `histogram` commands.

```
. graph box str, name(box_str)

. graph export box_str.png, name(box_str) replace
(file /Users/doylewr/lpo_prac/lessons/s1-07-cleaning/box_str.png written in PNG format)

. histogram str, name(hist_str)
(bin=20, start=14, width=1.6700001)

. graph export hist_str.png, name(hist_str) replace
(file /Users/doylewr/lpo_prac/lessons/s1-07-cleaning/hist_str.png written in PNG format)
/****
```

As you can see, there are two values of student teacher ratio that are big outliers. What should we do about these?

### Quick Exercise

There's another variable that has an extreme outlier. Find this variable and decide what to do about it.

### Impossible values

You should also look for impossible values. These include: negative values for things that must be positive like income or height; test scores that are above

the maximum; proportions that are negative or above one; or percentages that are negative or above 100.

Here's a summary of the `calw_pct` variable, which is expressed in percentage terms. It clearly has at least one impossible value.

```
. sum calw_pct
```

Variable	Obs	Mean	Std. Dev.	Min	Max
calw_pct	420	13.46038	12.25822	0	102.7

```
. replace calw_pct =. if calw_pct>100
(1 real change made, 1 to missing)
```

### Data that are off trend

When looking at data that are in panel format, very sharp changes from the previous period may be suspect. For instance, a student who goes from the 5th percentile to the 95th percentile in test scores.

### Checking relationships

Check to make sure that the variable is in the order that you would expect in the comparison. Are there high income students that are coded as low SES? Are there students with low GPAs and high test scores? These may be correct, but you need to check for strange patterns.

To test this, let's plot several of the variables against one another and look for problematic relationships. Here's a plot of average income, `avginc`, against the percent of students in the district on free or reduced price lunches, `meal_pct`:

```
. graph twoway scatter avginc meal_pct, name(sc_inc_meal)
```

```
. graph export sc_inc_meal.png, name(sc_inc_meal) replace
(file /Users/doylewr/lpo_prac/lessons/s1-07-cleaning/sc_inc_meal.png written in PNG format)
```

### Quick Exercise

There's another implausible value based on the relationship between two variables. Find it. (Hint: what is the biggest budget item in any school district?)

## Logically impossible combinations

Check that there aren't logically impossible combinations of variables. For example, we should be suspicious when a parent's age is less than that of the student. Sometimes your dataset documentation will alert you to potential problems. The codebook of the National Longitudinal Survey of Youth, 1997, helpfully explains that "researchers should note that [work hour] totals above 168 hours per week are suspect."

## Check calculations

Check to make sure that calculations have been done correctly. For example this dataset has several ratio measures. Let's recalculate these and see if they are correct. In the code below we recalculate the student teacher ratio and then plot it against the original calculation.

```
. gen str_two = enr_tot / teachers

. graph twoway scatter str_two str, name(sc_str_str_two)

. graph export sc_str_two.png, name(sc_str_str_two) replace
(file /Users/doylewr/lpo_prac/lessons/s1-07-cleaning/sc_str_two.png written in PNG format)
```

That's not good. Not good at all. `teachers` has a lot of the same -4 values. Maybe that means something. Clearly a school can't have -4 teachers. Let's recompute the ratio but only for schools with a positive number of teachers.

```
. replace teachers = . if teachers == -4
(67 real changes made, 67 to missing)

. gen str_three = enr_tot / teachers
(67 missing values generated)

. graph twoway scatter str_three str, name(sc_str_str_three)

. graph export sc_str_str_three.eps, name(sc_str_str_three) replace
(file sc_str_str_three.eps written in EPS format)
```

## Quick Exercise

There are other variables with problems in their calculations. Find them.

## Duplicates

You'll also need to look for duplicates in variables that shouldn't have any. The most obvious place to look is in *id* numbers or the equivalent. You should also check any other variable that ought to be unique. Luckily, Stata has a whole suite of duplicate commands to work with, including, the appropriately named `duplicates`.

Below is an example of a variable with no duplicates, followed by a variable with a couple of duplicates.

```
. duplicates report observation_number
```

Duplicates in terms of observation\_number

-----		
copies	observations	surplus
-----+-----		
1	420	0
-----		

```
. duplicates report dist_cod
```

Duplicates in terms of dist\_cod

-----		
copies	observations	surplus
-----+-----		
1	418	0
2	2	1
-----		

## Missing data

### Find missing data codes

First, figure out how missing data is coded from the codebook. Code those values as missing. That's a dot, `.`, in Stata.

```
. inspect testscr
```

testscr: API INDEX			-----			Number of Observations		
						Total	Integers	Nonintegers
	#	Negative	-	-	-	-	-	-
	#	Zero	-	-	-	-	-	-

	#	#		Positive	420	17	403
	#	#	#		-----	-----	-----
	#	#	#	Total	420	17	403
	#	#	#	Missing	-		
+-----							
605.55		706.75			420		
(More than 99 unique values)							

```
. inspect read_scr
```

read_scr:				Number of Observations			
-----				-----			
	#			Total	Integers	Nonintegers	
	#	Negative		77	77	-	
	#	Zero		15	15	-	
	#	Positive		328	33	295	
	#			-----	-----	-----	
	#	Total		420	125	295	
	#	Missing		-			
+-----							
-1		704		420			
(More than 99 unique values)							

```
. histogram read_scr, name(hist_read_scr)
(bin=20, start=-1, width=35.25)
```

```
. graph export hist_read_scr.png, name(hist_read_scr) replace
(file /Users/doylewr/lpo_prac/lessons/s1-07-cleaning/hist_read_scr.png written in PNG)
```

## Recode problematic data points as missing

Next, look for impossible values as above. These should also be recoded to missing.

## Dealing with zeros

Now, look for zeros. Are these really zeros? What was the criteria for having a zero? Check if these should really be missing.

Your best friend here is the Stata command **inspect**. Here's the result of the **inspect** command for two of the variables:

## Quick Exercise

Look for other variables with either negative or missing values and figure out what to do with them. Once you've got the data properly coded and missing entered as missing, use the `mvpatterns` command to figure out if there are systematic problems with your data.

## Conclusion

All of the above is only to get you started. There is no cookbook way to approach data cleaning. The idea is to get to know the data well enough that you know whether anomalies are just strange, but true, or problems in the data itself.

```
. log close
      name: <unnamed>
      log: /Users/doylewr/lpo_prac/lessons/s1-07-cleaning/cleaning.log
      log type: text
      closed on: 14 Oct 2020, 11:15:06
```

---

```
. exit
```